

JOHN DOE

San Francisco, CA · john.doe@email.com · rendercv.com · linkedin.com/in/rendercv · github.com/rendercv

WELCOME TO RENDERCV

RenderCV reads a CV written in a YAML file, and generates a PDF with professional typography.

Each section title is arbitrary.

You can choose any of the 9 entry types for each section.

Markdown syntax is supported everywhere. This is **bold**, *italic*, and link.

EDUCATION

PhD	Princeton University PhD in Computer Science — Princeton, NJ - Thesis: Efficient Neural Architecture Search for Resource-Constrained Deployment - Advisor: Prof. Sanjeev Arora - NSF Graduate Research Fellowship, Siebel Scholar (Class of 2022)	Sept 2018 – May 2023 4 years 9 months
BS	Boğaziçi University BS in Computer Engineering — Istanbul, Türkiye - GPA: 3.97/4.00, Valedictorian - Fulbright Scholarship recipient for Graduate Studies	Sept 2014 – June 2018 3 years 10 months

EXPERIENCE

Co-Founder & CTO , Nexus AI — San Francisco, CA - Built foundation model infrastructure serving 2M+ monthly API requests with 99.97% uptime - Raised \$18M Series A led by Sequoia Capital, with participation from a16z and Founders Fund - Scaled engineering team from 3 to 28 across ML research, platform, and applied AI divisions - Developed proprietary inference optimization reducing latency by 73% compared to baseline	June 2023 – present 3 years 1 month
Research Intern , NVIDIA Research — Santa Clara, CA - Designed sparse attention mechanism reducing transformer memory footprint by 4.2x - Co-authored paper accepted at NeurIPS 2022 (spotlight presentation, top 5% of submissions)	May 2022 – Aug 2022 4 months
Research Intern , Google DeepMind — London, UK - Developed reinforcement learning algorithms for multi-agent coordination - Published research at top-tier venues with significant academic impact - ICML 2022 main conference paper, cited 340+ times within two years - NeurIPS 2022 workshop paper on emergent communication protocols - Invited journal extension in JMLR (2023)	May 2021 – Aug 2021 4 months
Research Intern , Apple ML Research — Cupertino, CA - Created on-device neural network compression pipeline deployed across 50M+ devices - Filed 2 patents on efficient model quantization techniques for edge inference	May 2020 – Aug 2020 4 months
Research Intern , Microsoft Research — Redmond, WA - Implemented novel self-supervised learning framework for low-resource language modeling - Research integrated into Azure Cognitive Services, reducing training data requirements by 60%	May 2019 – Aug 2019 4 months

PROJECTS

FlashInfer Open-source library for high-performance LLM inference kernels - Achieved 2.8x speedup over baseline attention implementations on A100 GPUs - Adopted by 3 major AI labs, 8,500+ GitHub stars, 200+ contributors	Jan 2023 – present
NeuralPrune Automated neural network pruning toolkit with differentiable masks - Reduced model size by 90% with less than 1% accuracy degradation on ImageNet - Featured in PyTorch ecosystem tools, 4,200+ GitHub stars	Jan 2021

PUBLICATIONS

Sparse Mixture-of-Experts at Scale: Efficient Routing for Trillion-Parameter Models <i>John Doe, Sarah Williams, David Park</i> <i>NeurIPS 2023, July 2023</i>	July 2023
Neural Architecture Search via Differentiable Pruning <i>James Liu, John Doe</i> <i>NeurIPS 2022, Spotlight, Dec 2022</i>	Dec 2022
Multi-Agent Reinforcement Learning with Emergent Communication <i>Maria Garcia, John Doe, Tom Anderson</i> <i>ICML 2022, July 2022</i>	July 2022
On-Device Model Compression via Learned Quantization <i>John Doe, Kevin Wu</i> <i>ICLR 2021, Best Paper Award, May 2021</i>	May 2021

SELECTED HONORS

- MIT Technology Review 35 Under 35 Innovators (2024)
- Forbes 30 Under 30 in Enterprise Technology (2024)
- ACM Doctoral Dissertation Award Honorable Mention (2023)
- Google PhD Fellowship in Machine Learning (2020 – 2023)
- Fulbright Scholarship for Graduate Studies (2018)

SKILLS

Languages: Python, C++, CUDA, Rust, Julia
ML Frameworks: PyTorch, JAX, TensorFlow, Triton, ONNX
Infrastructure: Kubernetes, Ray, distributed training, AWS, GCP
Research Areas: Neural architecture search, model compression, efficient inference, multi-agent RL

PATENTS

1. Adaptive Quantization for Neural Network Inference on Edge Devices (US Patent 11,234,567)
2. Dynamic Sparsity Patterns for Efficient Transformer Attention (US Patent 11,345,678)
3. Hardware-Aware Neural Architecture Search Method (US Patent 11,456,789)

INVITED TALKS

4. Scaling Laws for Efficient Inference — Stanford HAI Symposium (2024)
3. Building AI Infrastructure for the Next Decade — TechCrunch Disrupt (2024)
2. From Research to Production: Lessons in ML Systems — NeurIPS Workshop (2023)
1. Efficient Deep Learning: A Practitioner’s Perspective — Google Tech Talk (2022)