

John Doe

📍 San Francisco, CA ✉ john.doe@email.com 🌐 rendercv.com 🌐 rendercv 🌐 rendercv

Welcome to RenderCV

RenderCV reads a CV written in a YAML file, and generates a PDF with professional typography.

Each section title is arbitrary.

You can choose any of the 9 entry types for each section.

Markdown syntax is supported everywhere. This is **bold**, *italic*, and [link](#).

Education

- Sept 2018 – May 2023 **Princeton University**, PhD in Computer Science – Princeton, NJ
- Thesis: Efficient Neural Architecture Search for Resource-Constrained Deployment
 - Advisor: Prof. Sanjeev Arora
 - NSF Graduate Research Fellowship, Siebel Scholar (Class of 2022)
- Sept 2014 – June 2018 **Boğaziçi University**, BS in Computer Engineering – Istanbul, Türkiye
- GPA: 3.97/4.00, Valedictorian
 - Fulbright Scholarship recipient for Graduate Studies

Experience

- June 2023 – present **Co-Founder & CTO**, Nexus AI – San Francisco, CA
- Built foundation model infrastructure serving 2M+ monthly API requests with 99.97% uptime
 - Raised \$18M Series A led by Sequoia Capital, with participation from a16z and Founders Fund
 - Scaled engineering team from 3 to 28 across ML research, platform, and applied AI divisions
 - Developed proprietary inference optimization reducing latency by 73% compared to baseline
- May 2022 – Aug 2022 **Research Intern**, NVIDIA Research – Santa Clara, CA
- Designed sparse attention mechanism reducing transformer memory footprint by 4.2x
 - Co-authored paper accepted at NeurIPS 2022 (spotlight presentation, top 5% of submissions)
- May 2021 – Aug 2021 **Research Intern**, Google DeepMind – London, UK
- Developed reinforcement learning algorithms for multi-agent coordination
 - Published research at top-tier venues with significant academic impact
 - ICML 2022 main conference paper, cited 340+ times within two years
 - NeurIPS 2022 workshop paper on emergent communication protocols
 - Invited journal extension in JMLR (2023)
- May 2020 – Aug 2020 **Research Intern**, Apple ML Research – Cupertino, CA
- Created on-device neural network compression pipeline deployed across 50M+ devices
 - Filed 2 patents on efficient model quantization techniques for edge inference

May 2019 – Aug 2019 **Research Intern**, Microsoft Research – Redmond, WA

- Implemented novel self-supervised learning framework for low-resource language modeling
- Research integrated into Azure Cognitive Services, reducing training data requirements by 60%

Projects

Jan 2023 – present **FlashInfer**

Open-source library for high-performance LLM inference kernels

- Achieved 2.8x speedup over baseline attention implementations on A100 GPUs
- Adopted by 3 major AI labs, 8,500+ GitHub stars, 200+ contributors

Jan 2021 **NeuralPrune**

Automated neural network pruning toolkit with differentiable masks

- Reduced model size by 90% with less than 1% accuracy degradation on ImageNet
- Featured in PyTorch ecosystem tools, 4,200+ GitHub stars

Publications

July 2023 **Sparse Mixture-of-Experts at Scale: Efficient Routing for Trillion-Parameter Models**

John Doe, Sarah Williams, David Park

[10.1234/neurips.2023.1234](https://arxiv.org/abs/10.1234/neurips.2023.1234) (NeurIPS 2023)

Dec 2022 **Neural Architecture Search via Differentiable Pruning**

James Liu, John Doe

[10.1234/neurips.2022.5678](https://arxiv.org/abs/10.1234/neurips.2022.5678) (NeurIPS 2022, Spotlight)

July 2022 **Multi-Agent Reinforcement Learning with Emergent Communication**

Maria Garcia, John Doe, Tom Anderson

[10.1234/icml.2022.9012](https://arxiv.org/abs/10.1234/icml.2022.9012) (ICML 2022)

May 2021 **On-Device Model Compression via Learned Quantization**

John Doe, Kevin Wu

[10.1234/iclr.2021.3456](https://arxiv.org/abs/10.1234/iclr.2021.3456) (ICLR 2021, Best Paper Award)

Selected Honors

- MIT Technology Review 35 Under 35 Innovators (2024)
- Forbes 30 Under 30 in Enterprise Technology (2024)
- ACM Doctoral Dissertation Award Honorable Mention (2023)
- Google PhD Fellowship in Machine Learning (2020 – 2023)
- Fulbright Scholarship for Graduate Studies (2018)

Skills

Languages: Python, C++, CUDA, Rust, Julia

ML Frameworks: PyTorch, JAX, TensorFlow, Triton, ONNX

Infrastructure: Kubernetes, Ray, distributed training, AWS, GCP

Research Areas: Neural architecture search, model compression, efficient inference, multi-agent RL

Patents

1. Adaptive Quantization for Neural Network Inference on Edge Devices (US Patent 11,234,567)
2. Dynamic Sparsity Patterns for Efficient Transformer Attention (US Patent 11,345,678)
3. Hardware-Aware Neural Architecture Search Method (US Patent 11,456,789)

Invited Talks

4. Scaling Laws for Efficient Inference – Stanford HAI Symposium (2024)
3. Building AI Infrastructure for the Next Decade – TechCrunch Disrupt (2024)
2. From Research to Production: Lessons in ML Systems – NeurIPS Workshop (2023)
1. Efficient Deep Learning: A Practitioner's Perspective – Google Tech Talk (2022)