

John Doe

📍 San Francisco, CA • ✉ john.doe@email.com • 🌐 rendercv.com • 🌐 rendercv • 🌐 rendercv

Welcome to RenderCV

RenderCV reads a CV written in a YAML file, and generates a PDF with professional typography.

Each section title is arbitrary.

You can choose any of the 9 entry types for each section.

Markdown syntax is supported everywhere. This is **bold**, *italic*, and [link](#).

Education

Princeton University – Princeton, NJ

Sept 2018 – May 2023

PhD in Computer Science

- Thesis: Efficient Neural Architecture Search for Resource-Constrained Deployment
- Advisor: Prof. Sanjeev Arora
- NSF Graduate Research Fellowship, Siebel Scholar (Class of 2022)

Boğaziçi University – Istanbul, Türkiye

Sept 2014 – June 2018

BS in Computer Engineering

- GPA: 3.97/4.00, Valedictorian
- Fulbright Scholarship recipient for Graduate Studies

Experience

Nexus AI, *Co-Founder & CTO* – San Francisco, CA

June 2023 – present

- Built foundation model infrastructure serving 2M+ monthly API requests with 99.97% uptime
- Raised \$18M Series A led by Sequoia Capital, with participation from a16z and Founders Fund
- Scaled engineering team from 3 to 28 across ML research, platform, and applied AI divisions
- Developed proprietary inference optimization reducing latency by 73% compared to baseline

NVIDIA Research, *Research Intern* – Santa Clara, CA

May 2022 – Aug 2022

- Designed sparse attention mechanism reducing transformer memory footprint by 4.2x
- Co-authored paper accepted at NeurIPS 2022 (spotlight presentation, top 5% of submissions)

Google DeepMind, *Research Intern* – London, UK

May 2021 – Aug 2021

- Developed reinforcement learning algorithms for multi-agent coordination
- Published research at top-tier venues with significant academic impact
- ICML 2022 main conference paper, cited 340+ times within two years
- NeurIPS 2022 workshop paper on emergent communication protocols
- Invited journal extension in JMLR (2023)

Apple ML Research, *Research Intern* – Cupertino, CA

May 2020 – Aug 2020

- Created on-device neural network compression pipeline deployed across 50M+ devices
- Filed 2 patents on efficient model quantization techniques for edge inference

Microsoft Research, Research Intern – Redmond, WA	May 2019 – Aug 2019
◦ Implemented novel self-supervised learning framework for low-resource language modeling ◦ Research integrated into Azure Cognitive Services, reducing training data requirements by 60%	

Projects	
FlashInfer	Jan 2023 – present
Open-source library for high-performance LLM inference kernels	
◦ Achieved 2.8x speedup over baseline attention implementations on A100 GPUs ◦ Adopted by 3 major AI labs, 8,500+ GitHub stars, 200+ contributors	
NeuralPrune	Jan 2021
Automated neural network pruning toolkit with differentiable masks	
◦ Reduced model size by 90% with less than 1% accuracy degradation on ImageNet ◦ Featured in PyTorch ecosystem tools, 4,200+ GitHub stars	

Publications	
Sparse Mixture-of-Experts at Scale: Efficient Routing for Trillion-Parameter Models	July 2023
John Doe, Sarah Williams, David Park	
10.1234/neurips.2023.1234 (NeurIPS 2023)	
Neural Architecture Search via Differentiable Pruning	Dec 2022
James Liu, John Doe	
10.1234/neurips.2022.5678 (NeurIPS 2022, Spotlight)	
Multi-Agent Reinforcement Learning with Emergent Communication	July 2022
Maria Garcia, John Doe, Tom Anderson	
10.1234/icml.2022.9012 (ICML 2022)	
On-Device Model Compression via Learned Quantization	May 2021
John Doe, Kevin Wu	
10.1234/iclr.2021.3456 (ICLR 2021, Best Paper Award)	

Selected Honors	
◦ MIT Technology Review 35 Under 35 Innovators (2024) ◦ Forbes 30 Under 30 in Enterprise Technology (2024) ◦ ACM Doctoral Dissertation Award Honorable Mention (2023) ◦ Google PhD Fellowship in Machine Learning (2020 – 2023) ◦ Fulbright Scholarship for Graduate Studies (2018)	

Skills	
Languages: Python, C++, CUDA, Rust, Julia	
ML Frameworks: PyTorch, JAX, TensorFlow, Triton, ONNX	
Infrastructure: Kubernetes, Ray, distributed training, AWS, GCP	
Research Areas: Neural architecture search, model compression, efficient inference, multi-agent RL	

Patents	
1. Adaptive Quantization for Neural Network Inference on Edge Devices (US Patent 11,234,567)	
2. Dynamic Sparsity Patterns for Efficient Transformer Attention (US Patent 11,345,678)	
3. Hardware-Aware Neural Architecture Search Method (US Patent 11,456,789)	

Invited Talks

4. Scaling Laws for Efficient Inference – Stanford HAI Symposium (2024)
3. Building AI Infrastructure for the Next Decade – TechCrunch Disrupt (2024)
2. From Research to Production: Lessons in ML Systems – NeurIPS Workshop (2023)
1. Efficient Deep Learning: A Practitioner’s Perspective – Google Tech Talk (2022)